

- Name: Prerna Sunil Jadhav
- Sap ID: 60004220127
- Batch: C22
- Branch: Computer Engineering
- Course: Machine Learning
- Experiment 3: Logistics Regression

## USING LIBRARY

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

data=pd.read_csv("../content/framingham.csv")
data.head()
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4238 entries, 0 to 4237
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   male                  4238 non-null   int64  
 1   age                   4238 non-null   int64  
 2   education             4133 non-null   float64 
 3   currentSmoker         4238 non-null   int64  
 4   cigsPerDay            4209 non-null   float64 
 5   BPMeds                4185 non-null   float64 
 6   prevalentStroke       4238 non-null   int64  
 7   prevalentHyp          4238 non-null   int64  
 8   diabetes              4238 non-null   int64  
 9   totChol               4188 non-null   float64 
10   sysBP                 4238 non-null   float64 
11   diaBP                 4238 non-null   float64 
12   BMI                   4219 non-null   float64 
13   heartRate             4237 non-null   float64 
14   glucose               3850 non-null   float64 
15   TenYearCHD            4238 non-null   int64  
dtypes: float64(9), int64(7)
memory usage: 529.9 KB
```

```
sns.countplot(data,x='diabetes')
```

```
data =data.dropna()
```

```
data.isnull().sum()
```

```
male          0
age           0
education     0
currentSmoker 0
cigsPerDay    0
BPMeds        0
prevalentStroke 0
prevalentHyp  0
diabetes      0
totChol       0
sysBP         0
diaBP         0
BMI           0
heartRate     0
glucose       0
TenYearCHD    0
dtype: int64
```

```
data.shape
```

```
(3656, 16)
```

```
x=data.iloc[:, :8]
y=data.iloc[:, 8]
y.head()
```

```
0    0
1    0
2    0
3    0
4    0
Name: diabetes, dtype: int64
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)
```

```
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
new_xtrain=sc.fit_transform(x_train)
new_xtest=sc.transform(x_test)
```

```
from sklearn.linear_model import LogisticRegression
classifier=LogisticRegression()
classifier.fit(new_xtrain,y_train)
```

```
y_pred=classifier.predict(new_xtest)
```

```
y_pred
```

```
array([0, 0, 0, ..., 0, 0, 0])
```

```
from sklearn.metrics import confusion_matrix
confusion_matrix(y_test, y_pred)
```

```
array([[1070,    0],
       [   27,    0]])
```

```
from sklearn.metrics import accuracy_score
print('Accuracy = ' , accuracy_score(y_test, y_pred))
```

```
Accuracy = 0.97538742023701
```

## WITHOUT USING LIBRARIES

```
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn import datasets
import matplotlib.pyplot as plt
```

```
d = datasets.load_breast_cancer()
x, y = d.data, d.target
```

```
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.2, random_state=1234)
```

```
class LogisticRegression:

    def __init__(self, learning_rate=0.001, n_iters=1000):
        self.lr = learning_rate
        self.n_iters = n_iters
        self.weights = None
        self.bias = None

    def fit(self, X, y):
        n_samples, n_features = X.shape

        # init parameters
        self.weights = np.zeros(n_features)
        self.bias = 0

        # gradient descent
        for _ in range(self.n_iters):

            linear_model = np.dot(X, self.weights) + self.bias
            # apply sigmoid function
            y_predicted = self._sigmoid(linear_model)

            # compute gradients
            dw = (1 / n_samples) * np.dot(X.T, (y_predicted - y)) #derivative w.r.t weights
            db = (1 / n_samples) * np.sum(y_predicted - y) #derivative w.r.t bias
            # update parameters
            self.weights -= self.lr * dw
            self.bias -= self.lr * db

    def predict(self, X):
        linear_model = np.dot(X, self.weights) + self.bias
        y_predicted = self._sigmoid(linear_model)
        y_predicted_cls = [1 if i > 0.5 else 0 for i in y_predicted]
        print(y_predicted_cls)
        return np.array(y_predicted_cls)

    def _sigmoid(self, x):
        return 1 / (1 + np.exp(-x))

def accuracy(y_true, y_pred):
    accuracy = np.sum(y_true == y_pred) / len(y_true)
    return accuracy
```

```
itr=[]
acc=[]
```

```
regressor = LogisticRegression(learning_rate=0.0001, n_iters=1000)
regressor.fit(xtrain, ytrain)
predictions = regressor.predict(xtest)
itr.append(1000)
```

[1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1,

```
print("Logistic Regression classification accuracy:", accuracy(ytest, predictions))
acc.append(accuracy(ytest, predictions))
```

Logistic Regression classification accuracy: 0.9298245614035088

```
regressor.weights
```

```
array([ 3.15267538e-02,  4.38592690e-02,  1.82394637e-01,  7.27657289e-02,
        2.81683690e-04, -1.58921860e-04, -5.94869592e-04, -2.47270611e-04,
        5.51932783e-04,  2.26761495e-04,  1.51071202e-04,  3.05608006e-03,
       -1.13197589e-03, -8.16912730e-02,  1.44483806e-05, -4.65383514e-05,
       -7.39608956e-05, -1.01074526e-05,  5.54703739e-05,  1.84514942e-06,
        3.33636766e-02,  5.58146134e-02,  1.82897269e-01, -9.88462272e-02,
        3.48754542e-04, -6.74411296e-04, -1.27433782e-03, -2.94192307e-04,
        7.76822709e-04,  1.96607498e-04])
```

```
regressor.bias
```

```
0.004111914763563403
```